AI-assisted marking for course discussions

A human-led, evidence-linked strategy for participation, quality and development over time

Core proposition Chisimba uses AI to prepare an inspectable marking draft—not to make or publish the final academic judgement. Every suggested criterion score points back to the learner's actual discussion contributions, and a lecturer must review and release the result.

Prepared for educators, psychologists, learning designers, markers and institutional leaders considering responsible AI-supported assessment.

1. Why discussions deserve a richer assessment model

Online discussion is not merely a count of posts. It can reveal how learners reason in public, respond to other perspectives, use evidence, sustain participation and develop ideas across time. Conventional analytics often collapse this into activity totals; conventional manual marking can become prohibitively slow. Chisimba's strategy combines transparent computational assistance with accountable human judgement.

- Participation matters, but frequency alone is not quality.
- A polished isolated post should not automatically outweigh sustained, constructive engagement.
- Improvement can be meaningful evidence, but it must never become an artificial requirement.
- Learners who perform strongly from the outset must be able to earn the highest level without displaying a manufactured growth curve.

2. The assessment design

A course Discussion can be configured as an ordinary participation activity or as a marked activity. When marked, it becomes eligible for the course Assessment Plan and Gradebook. Discussion owns the conversations, rubric evidence and authoritative marks; Gradebook owns plan inclusion, weighting and contribution to the course result.

Design principle The same activity remains usable without AI. If no AI service is configured, or a suggestion fails, lecturers use the identical rubric and marking workspace manually.

3. Built-in participation and quality rubric

The protected default rubric totals 100 points. Lecturers can later select or adapt an appropriate course rubric, but a defensible default is available immediately.

Criterion	Weight	What strong evidence looks like
Participation, consistency and development	25%	Timely, purposeful contributions sustained at a high standard, or meaningful development toward that standard.
Relevance to the discussion	25%	Focused contributions that respond to the particular conversation and help it move forward.
Quality of reasoning and evidence	30%	Accurate, supported and critically reasoned contributions rather than unsupported assertion.
Constructive engagement and responsiveness	20%	Respectful engagement that builds on perspectives, invites participation, or thoughtfully applies feedback.

4. Improvement over time: allowed, not required

The rubric treats development as one legitimate form of strong evidence. It does not assume that every learner begins from the same point, follows a linear trajectory, or must initially perform poorly to demonstrate learning.

- Consistently excellent work can satisfy the highest descriptor.
- Genuine improvement or thoughtful response to feedback can also strengthen a judgement.
- Mere change is not improvement and receives no automatic reward.
- A temporary dip, experimentation or uneven participation is interpreted in context rather than forced into a simple upward trend.
- AI instructions explicitly prohibit penalising stable high performance or inventing a growth narrative.

5. What the AI actually does

1. Chisimba freezes a bounded evidence set containing the learner's contributions to the selected Discussion, with post identifiers, dates, topics and plain text.
2. The AI receives the Discussion purpose, the versioned rubric and only that learner's permitted evidence.
3. It suggests a bounded score for every criterion, a rationale, supporting post identifiers and draft formative feedback.

4. It flags insufficient evidence instead of filling gaps with assumptions.
5. The lecturer sees an editable draft alongside the evidence, changes anything necessary and explicitly saves the reviewed mark.
6. Only the lecturer-reviewed result becomes authoritative and available to Gradebook.

6. Human authority and safeguards

Safeguard	Implementation logic
No autonomous grading	AI output is a draft. It cannot publish a mark or update Gradebook directly.
Evidence traceability	Every rationale must cite canonical post IDs; invented or out-of-scope IDs are rejected.
Frozen audit evidence	The evidence snapshot used for a job is retained so later edits cannot rewrite the historical basis of the suggestion.
Bounded scores	Criterion scores are clamped to the rubric maximum on the server.
Scope and role control	Only authorised course lecturers can request suggestions or save marks for learners in that course.
No authorship detection	The AI may not classify writing as AI- or human-authored, infer identity, personality, intent or protected characteristics.
Failure-safe design	Missing AI, incomplete output or insufficient evidence returns a clear non-marking state; manual marking remains available.
Replay protection	Marking and AI requests use one-time, action- and learner-specific security tokens.
Versioned provenance	Rubric version, evidence, requester, timestamps, result and error state are retained for audit.

7. Psychological and educational considerations

Construct validity

The system separates observable participation from the quality of reasoning and interpersonal engagement. This reduces the risk that a simple behavioural proxy—such as post count—is mistaken for learning.

Fairness and accessibility

Participation patterns can reflect language, disability, connectivity, confidence, culture, caring responsibilities and course design. The rubric therefore evaluates submitted evidence without inferring personal causes. Lecturers retain responsibility for accommodations, missing context and alternative evidence.

Automation bias

Criterion-level evidence and editable rationales are deliberately visible to slow uncritical acceptance. The interface describes the result as a suggestion, not a recommendation that carries institutional authority.

Behavioural effects

Learners should know what is assessed and how. The rubric should encourage substantive interaction, not strategic posting, artificial agreement or performative frequency. Course guidance should state that concise, valuable contributions may be stronger than numerous low-value posts.

Data minimisation

Only contributions belonging to the selected learner and Discussion are supplied. Peer posts are not required merely to score the learner, although quoted context may need careful future treatment. Evidence is bounded in count and length.

8. Lecturer journey

1. Create or edit a course Discussion and enable “Available as a course activity.”
2. Optionally enable “Marked discussion,” choose formative or summative classification and set the total mark.
3. Add the Discussion to a chapter from the Assessments palette, identified by the Discussion icon rather than a generic clipboard.
4. Open the marking workspace to see learners, contribution counts, existing marks and the complete rubric.
5. Mark manually, or request an AI suggestion for an individual learner.
6. Review every criterion, inspect cited posts, edit scores and feedback, then save the reviewed result.
7. Include the marked Discussion in the Assessment Plan and assign its course weighting. Gradebook consumes the reviewed percentage through the standard provider API.

9. Technical architecture in plain language

Chisimba consumes its own APIs and provider contracts. Discussion exposes two deliberate discovery views: the course-content palette can see all Discussions explicitly enabled as course activities, while Gradebook sees only those also enabled for marking. This prevents an unmarked conversation from accidentally becoming a Gradebook item.

- Discussion provider: lists eligible activities, validates course scope, launches the activity and returns reviewed learner results.
- Context Content: places a reference to the Discussion in a chapter without copying its posts.
- Gradebook: owns the assessment plan, short label, ordering, weight and contribution calculation.
- AI job queue: performs bounded background work and retains evidence-linked drafts and failures.
- Rubric service: provisions a protected, versioned default while allowing appropriate future course-specific rubrics.

10. Suggested video narrative

A concise presentation can follow this sequence:

1. The problem: discussion assessment is valuable but difficult to mark consistently at scale.
2. The course setup: enable a Discussion as an activity and optionally as marked work.
3. The learner experience: authentic conversation remains the centre of the activity.
4. The rubric: participation, relevance, reasoning and constructive engagement—including optional evidence of development.
5. The AI demonstration: request a draft and reveal criterion scores, rationales and exact evidence links.
6. The critical safeguard: change a suggested score and explicitly release the lecturer's judgement.
7. The architecture: the same provider powers chapter placement, Assessment Plan inclusion and Gradebook results.
8. The message: Chisimba is using AI to make academic judgement more inspectable and scalable, not to remove the academic.

11. Claims we can responsibly make

Presentation language Chisimba offers human-led, rubric-based AI assistance for discussion marking. Suggestions are criterion-level, evidence-linked, bounded, auditable and never published without lecturer review.

Avoid claiming that the system eliminates bias, detects AI authorship, objectively measures learning, or makes lecturers unnecessary. Its value is narrower and stronger: it provides structured assistance, visible evidence and accountable human control.

12. Acceptance standard

- A Discussion can be an unmarked or marked course activity.
- The Assessments palette uses the skin’s Discussion icon.
- Only marked Discussions are discoverable by Gradebook.
- Manual marking works without an AI service.
- The default rubric is available immediately and totals 100 points.
- Improvement can support a judgement but is never required.
- AI suggestions cite retained evidence and cannot write authoritative marks.
- Lecturer review, scope protection, replay protection and audit records are enforced server-side.